

Presentation of Leads Case Study

Furthermore, I have described the process that I have followed to do my Leads case study.

Problem Statement

- X Education is a company that markets courses online through websites and search engines. It's lead conversion rate is quite poor.
- In this case study we try to create a logistic regression machine learning model from the data provided by the education company of customers who have been contacted for courses in order to understand what data may play an important role in increasing the lead conversion rate so that the company can focus more on communicating the potential leads that may lean towards choosing a course.

First step upload the required libraries

- In this step, I have uploaded the required libraries like a panda, seaborn, matplotlib, sklearn, statsmodel, etc.
- This you can see in the diagram, I have attached to the left.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score
```

Import the CSV data to the notebook

- In this step, I have used the panda library to import the CSV data to the Jupiter notebook by using the `read_csv` function.
- After that I have used `leads.head()` to see the first few lines of the data that I have been imported.

```
In [2]: leads = pd.read_csv("Downloads/Leads.csv")
In [3]: leads.head()
```

Out[3]:

| | Prospect ID | Lead Number | Lead Origin | Lead Source | Do Not Email | Do Not Call | Converted | TotalVisits | Total Time Spent on Website | Page Views Per Visit | Get updates on DM Content | Lead Profile | City | Asymmetrique Activity Index | Asymm Profile |
|---|--------------------------------------|-------------|-------------------------|----------------|--------------|-------------|-----------|-------------|-----------------------------|----------------------|---------------------------|--------------|----------------|-----------------------------|---------------|
| 0 | 7927b2df-8bba-4d29-b9a2-b6e0beafe620 | 660737 | API | Olark Chat | No | No | 0 | 0.0 | 0 | 0.0 | ... | No | Select | Select | 02.Medium |
| 1 | 2a272436-5132-4136-86fa-dcc88c88f482 | 660728 | API | Organic Search | No | No | 0 | 5.0 | 674 | 2.5 | ... | No | Select | Select | 02.Medium |
| 2 | 8cc8c611-a219-4f35-ad23-fdfd2656bd8a | 660727 | Landing Page Submission | Direct Traffic | No | No | 1 | 2.0 | 1532 | 2.0 | ... | No | Potential Lead | Mumbai | 02.Medium |

Perform some operations

- Here I performed some operations to see the data in-depth like `inp0.shape` (to see the number of rows and columns), `inp0.dtypes`(to see the data types of the different columns)
- `inp0.info()` to see the information of the columns.

```
In [4]: M leads.shape
```

```
Out[4]: (9240, 37)
```

```
In [5]: M leads.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9240 entries, 0 to 9239
Data columns (total 37 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Prospect ID                             9240 non-null   object
1   Lead Number                             9240 non-null   int64
2   Lead Origin                             9240 non-null   object
3   Lead Source                             9204 non-null   object
4   Do Not Email                            9240 non-null   object
5   Do Not Call                             9240 non-null   object
6   Converted                               9240 non-null   int64
7   TotalVisits                             9103 non-null   float64
8   Total Time Spent on Website              9240 non-null   int64
9   Page Views Per Visit                     9103 non-null   float64
10  Last Activity                            9137 non-null   object
11  Country                                  6779 non-null   object
12  Specialization                           7802 non-null   object
```

Apply Data Cleaning on the lead data set

- The very First step is to see how many null values the columns have and calculate their percentage by using `is_null().sum()`.
- Second, drop those columns which having null values greater than 45% by using the `drop` function.

```
In [12]: # checking the null values in terms of percentage
         round(100*(leads.isnull().sum()/len(leads.index)),2)

Out[12]: Prospect ID          0.00
         Lead Number         0.00
         Lead Origin         0.00
         Lead Source         0.39
         Do Not Email        0.00
         Do Not Call         0.00
         Converted           0.00
         TotalVisits         1.48
         Total Time Spent on Website 0.00
         Page Views Per Visit 1.48
         Last Activity        1.11
         Country             26.63
         Specialization       36.58
         How did you hear about X Education 78.46
         What is your current occupation 29.11
         What matters most to you in choosing a course 29.32
         Search              0.00
         Magazine            0.00
         Newspaper Article    0.00
         X Education Forums   0.00
         Newspaper           0.00
```

Data Cleaning

- Next, the remaining columns which are having few null values. Replacing the null values with the **most frequent value**.
- There are some columns which have unbalanced data. So I removed those columns such as **Do Not Call, Search, Magazine**, etc.

```
In [68]: leads['City'].value_counts(dropna=False)
```

```
Out[68]: NaN                3669  
Mumbai                3222  
Thane & Outskirts      752  
Other Cities           686  
Other Cities of Maharashtra  457  
Other Metro Cities     380  
Tier II Cities         74  
Name: City, dtype: int64
```

```
In [69]: # replacing Nan values with most occur value i.e mumbai  
leads['City'] = leads['City'].replace(np.nan, 'Mumbai')
```


```
In [70]: leads['City'].value_counts(dropna=False)
```

```
Out[70]: Mumbai                6891  
Thane & Outskirts      752  
Other Cities           686  
Other Cities of Maharashtra  457  
Other Metro Cities     380  
Tier II Cities         74  
Name: City, dtype: int64
```

Data Cleaning

- As a result of the previous operation, I don't have any null values in my dataset and the data is balanced.
- Next, I handle the outliers for the numeric data such as **TotalVisits, Total Time Spent on Website and Page Views Per Visit.**

```
In [112]: #Total Visits  
plt.figure(figsize=(10,5))  
sns.boxplot(y=leads['TotalVisits'])  
plt.show()
```



Creating Dummy variables

- Creating the dummy variables for the categorical columns which are having more than two categories.
- After creating the dummy variables, I left with 103 columns and the data is ready for the machine learning algorithm.

```
##### Lead Origin, Lead Source #####

In [133]: # getting dummies and adding the result to the Leads
dummy = pd.get_dummies(leads[['Lead Origin', 'Lead Source', 'Last Activity',
                              'Specialization', 'What is your current occupation', 'Tags', 'City',
                              'Last Notable Activity']], drop_first=True)

leads = pd.concat([leads,dummy],1)

In [134]: leads.head()
```

Out[134]:

Splitting the data in to train and test

- Next step, is to split data into train and test by using the **train-test library**.
- Scaling the numeric variable by using the **StandardScaler method**.

```
In [139]: # Splitting the data into train and test
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.7, random_state=100)
```

```
In [140]: X_train.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 6468 entries, 1871 to 5640
Columns: 102 entries, TotalVisits to Last Notable Activity_View in browser link Clicked
dtypes: float64(2), int64(1), uint8(99)
memory usage: 827.4 KB
```

```
In [141]: #scaling numeric columns

scaler = StandardScaler()

numeric_cols = X_train.select_dtypes(include=['float64', 'int64']).columns

X_train[numeric_cols] = scaler.fit_transform(X_train[numeric_cols])

X_train.head()
```

Creating the models

- After scaling the data, we create models using RFE until we get the p-values of all the columns less than 0.05.
- We also check the VIF if it is less than 3 or not.
- We eliminate the columns which are having high p-values and VIF one by one.
- When we find the model suitable for the analysis we continue with the process and do the y prediction.
- We predict the value across the target variable and find its accuracy, precision, recall and specificity.



| | | | |
|------------------|------------------|-------------------|----------|
| Dep. Variable: | Converted | No. Observations: | 6468 |
| Model: | GLM | Df Residuals: | 6452 |
| Model Family: | Binomial | Df Model: | 15 |
| Link Function: | logit | Scale: | 1.0000 |
| Method: | IRLS | Log-Likelihood: | -1396.4 |
| Date: | Sun, 16 Oct 2022 | Deviance: | 2792.7 |
| Time: | 17:44:00 | Pearson chi2: | 1.06e+04 |
| No. Iterations: | 8 | | |
| Covariance Type: | nonrobust | | |

| | coef | std err | z | P> z | [0.025 | 0.975] |
|--|---------|---------|---------|-------|--------|--------|
| const | -2.5631 | 0.088 | -29.114 | 0.000 | -2.736 | -2.391 |
| What is your current occupation_Unemployed | 2.2634 | 0.117 | 19.370 | 0.000 | 2.034 | 2.492 |
| What is your current occupation_Working Professional | 2.5416 | 0.356 | 7.149 | 0.000 | 1.845 | 3.238 |
| Lead Source_Welingak Website | 2.9473 | 0.733 | 4.021 | 0.000 | 1.511 | 4.384 |
| Last Activity_SMS Sent | 2.0669 | 0.109 | 18.950 | 0.000 | 1.853 | 2.281 |
| Tags_Already a student | -4.8662 | 0.718 | -6.774 | 0.000 | -6.274 | -3.458 |
| Tags_Closed by Horizzon | 5.8904 | 1.010 | 5.829 | 0.000 | 3.910 | 7.871 |
| Tags_Graduation in progress | -2.5781 | 0.493 | -5.228 | 0.000 | -3.545 | -1.612 |
| Tags_Interested in full time MBA | -3.8136 | 0.736 | -5.184 | 0.000 | -5.255 | -2.372 |
| Tags_Interested in other courses | -3.4071 | 0.325 | -10.479 | 0.000 | -4.044 | -2.770 |
| Tags_Lost to EINS | 5.5233 | 0.727 | 7.601 | 0.000 | 4.099 | 6.948 |
| Tags_Not doing further education | -4.6677 | 1.015 | -4.599 | 0.000 | -6.657 | -2.678 |
| Tags_Other_tags | -3.0611 | 0.271 | -11.300 | 0.000 | -3.592 | -2.530 |
| Tags_Ringing | -4.4302 | 0.233 | -19.033 | 0.000 | -4.886 | -3.974 |
| Tags_Will revert after reading the email | 3.3450 | 0.183 | 18.256 | 0.000 | 2.986 | 3.704 |

| | Features | VIF |
|----|---|------|
| 5 | Tags_Closed by Horizzon | 1.33 |
| 11 | Tags_Other_tags | 1.29 |
| 14 | Tags_switched off | 1.24 |
| 10 | Tags_Not doing further education | 1.14 |
| 7 | Tags_Interested in full time MBA | 1.10 |
| 2 | Lead Source_Welingak Website | 1.09 |
| 6 | Tags_Graduation in progress | 1.09 |
| 9 | Tags_Lost to EINS | 1.06 |
| 1 | What is your current occupation_Working Profes... | 0.95 |
| 8 | Tags_Interested in other courses | 0.44 |
| 13 | Tags_Will revert after reading the email | 0.32 |
| 4 | Tags_Already a student | 0.29 |
| 12 | Tags_Ringing | 0.20 |
| 3 | Last Activity_SMS Sent | 0.12 |
| 0 | What is your current occupation_Unemployed | 0.11 |

ROC Curve

- Before the ROC Curve, we are using a 0.5 cutoff for all the observations. But the ROC curve gives us the best cutoff value.
- We use this cutoff value to predict the output across the target variable.
- We again find accuracy, precision, and specificity which are the final solution for the train data.
- Next step is to apply the model to the test data which is not seen by the model till now.
- Depending on the outcome we find out how the model is behaving on the test data set.



Results and Conclusion

- From the final ML model we get 91% accuracy, sensitivity of around 81% and specificity of around 97% and after selecting the best cut-off which is 0.3 we calculate the accuracy of 90%, sensitivity of 94% and specificity of around 88% and after applying model to the test data set we concluded that model gives us accuracy of 90%, sensitivity of 96%, specificity of around 87%, precision score of 83% and recall score of 96%.
- From these results as determine that the model has made good predictions and the company should focus on leads that are closed by 'Horizzon', lost to EINS, should avoid leads that are too busy/marked as switched off, already a student and who are not interested in further education. The marketing team should also consider leads that are interested in full time MBA courses and will revert back after reading the email.